

SI252 Reinforcement Learning: Homework #06

Due on June 1, 2025 at 11:59 p.m.(CST)

Name: **Zhou Shouchen**
Student ID: 2021533042

Problem 1

Required: OpenAI Spinning Up in Deep RL. Welcome to Spinning Up in Deep RL! This is an educational resource produced by OpenAI that makes it easier to learn about deep reinforcement learning (deep RL). Please study the documents and install the environment with PyTorch version.

- (a) Finish the problem set 1: “Basics of Implementation”. It includes three exercises: Gaussian Log-Likelihood, Policy for PPO, and Computation Graph for TD3.
- (b) Finish the problem set 2: “Algorithm Failure Modes”. It includes two exercises: Value Function Fitting in TRPO, and Silent Bug in DDPG.

Solution

(a) After setting up the environment in the methods in ‘README.md’, we can successfully run the codes. The codes could be check in ‘code/spinningup/spinup/exercises/pytorch/problem_set_1’.

- exercise 1_1: For a k dimension diagonal Gaussian distribution $\mathbf{x} \sim \mathcal{N}(\boldsymbol{\mu}, \boldsymbol{\Sigma})$, where $\boldsymbol{\Sigma} = \text{diag}(\sigma_1^2, \dots, \sigma_k^2)$, thus its PDF is

$$p(\mathbf{x}; \boldsymbol{\mu}, \boldsymbol{\Sigma}) = \frac{1}{(2\pi)^{\frac{k}{2}} |\boldsymbol{\Sigma}|^{\frac{1}{2}}} \exp \left[-\frac{1}{2} (\mathbf{x} - \boldsymbol{\mu})^\top \boldsymbol{\Sigma}^{-1} (\mathbf{x} - \boldsymbol{\mu}) \right]$$

Where

$$|\boldsymbol{\Sigma}| = \prod_{i=1}^k \sigma_i^2, \quad \boldsymbol{\Sigma}^{-1} = \text{diag}(\sigma_1^{-2}, \dots, \sigma_k^{-2}), \quad (\mathbf{x} - \boldsymbol{\mu})^\top \boldsymbol{\Sigma}^{-1} (\mathbf{x} - \boldsymbol{\mu}) = \sum_{i=1}^k \frac{(x_i - \mu_i)^2}{\sigma_i^2}$$

So above all, the log-likelihood is

$$\begin{aligned} \log p(\mathbf{x}; \boldsymbol{\mu}, \boldsymbol{\Sigma}) &= -\frac{k}{2} \log(2\pi) - \sum_{i=1}^k \log \sigma_i - \frac{1}{2} \sum_{i=1}^k \frac{(x_i - \mu_i)^2}{\sigma_i^2} \\ &= -\frac{1}{2} \sum_{i=1}^k \left(\log(2\pi) + 2 \log \sigma_i + \left(\frac{x_i - \mu_i}{\sigma_i} \right)^2 \right) \end{aligned}$$

The implementation are check results are as follows:

The screenshot shows a Jupyter Notebook interface with a code editor and a terminal. The code editor contains the following Python code:

```

15
16 def gaussian_likelihood(x, mu, log_std):
17     """
18     Args:
19         x: Tensor with shape [batch, dim]
20         mu: Tensor with shape [batch, dim]
21         log_std: Tensor with shape [batch, dim] or [dim]
22     Returns:
23         Tensor with shape [batch]
24     """
25     #####
26     # YOUR CODE HERE #
27     #####
28     pre_sum = -0.5 * (np.log(2 * np.pi) + 2 * log_std + ((x - mu) / torch.exp(log_std)) ** 2)
29     return pre_sum.sum(axis=-1)

```

The terminal shows the command to run the code:

```

• (spinningup) cellverse@cellverse:/media/cellverse/share013/zhoushch/ShanghaiTech-SI252/homework/hw6/code/spinningup$ python -m spinup.exercises.pytorch.problem_set_1.exercise1_1

```

The output of the terminal shows the message:

```

Congratulations! Your answer is correct.

```

- exercise 1_2: The implementation are check results are as follows:

Epoch	19
AverageEpRet	1e+03
StdEpRet	0
MaxEpRet	1e+03
MinEpRet	1e+03
EpLen	1e+03
AverageVVals	99.1
StdVVals	0.976
MaxVVals	100
MinVVals	92.1
TotalEnvInteracts	8e+04
LossPi	5.96e-11
LossV	1.64
DeltaLossPi	-0.00734
DeltaLossV	-0.703
Entropy	0.717
KL	0.00357
ClipFrac	0.0385
StopIter	79
Time	8.6


=====


Congratulations! Your answer is correct.

=====

- exercise 1_3: According to the discription, within 10 rounds, the score in HalfSheetah should exceed 300, while the score in InvertedPendulum should reach 150.

And the following curves show that the performance achieved the expected results:






(b) Run the commands in the 'README.md', we can get the following results. The codes could be check in 'code/spinningup/spinup/exercises/pytorch/problem_set_2'.


- exercise 2_1:

The difference is quite substantial: with a trained value function, the agent is able to quickly make progress. With an untrained value function, the agent gets stuck early on.


The curve of v0 with different seeds are as follows:




The curve of v80 with different seeds are as follows:



The comparison between curve of v0 and v80, and their average and variance are as follows:




- exercise 2_2: The curve of the code without bug and with different seeds are as follows:



The curve of the code with bug and with different seeds are as follows:



The comparison between curve with and without bug, and their average and variance are as follows:



The code's bug is shown in the following figure:

```

36 class BuggedMLPQFunction(nn.Module):
37
38     def __init__(self, obs_dim, act_dim, hidden_sizes, activation):
39         super().__init__()
40         self.q = mlp([obs_dim + act_dim] + list(hidden_sizes) + [1], activation)
41
42     def forward(self, obs, act):
43         q = self.q(torch.cat([obs, act], dim=-1))
44         q_squeezed = torch.squeeze(q, -1)
45         # print(q.shape) ([64, 1])
46         # print(q_squeezed.shape) ([64])
47         return self.q(torch.cat([obs, act], dim=-1))
48

```

The error in the code is due to a missing `squeeze(-1)` operation, which results in the output tensor having a shape of `[B, 1]` instead of the expected `[B]`. This can cause issues in subsequent operations, as many functions (such as loss computation or tensor arithmetic) expect inputs with shape `[B]`, and using `[B, 1]` may lead to shape mismatches or broadcasting errors.

- Bonus: Are there any choices of hyperparameters which would have hidden the effects of the bug? Adjusting the hyperparameter: set `batchsize` into 1, in this case, the bug disappears.

Problem 2

Optional: Hugging Face Course in Deep RL. Welcome to Hugging Face Course in Deep RL! This is an educational resource produced by Hugging Face community that makes it easier to learn about deep reinforcement learning (deep RL).

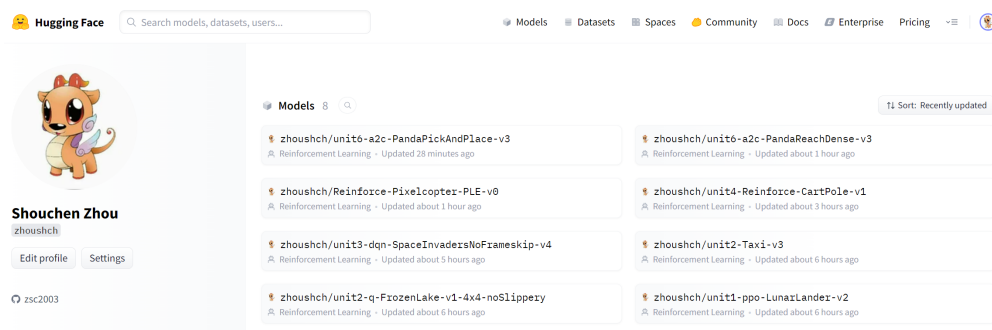
(a) Finish basic units including unit 1(Introduction to Deep Reinforcement Learning), unit 2(Introduction to Q-Learning), unit 3(Deep Q-Learning with Atari Games), unit 4(Policy Gradient with Pytorch), unit 6(Actor-Critic Methods with Robotics Environments), unit 8(Proximal Policy Optimization)

(b) (Bonus) Finish all remaining units

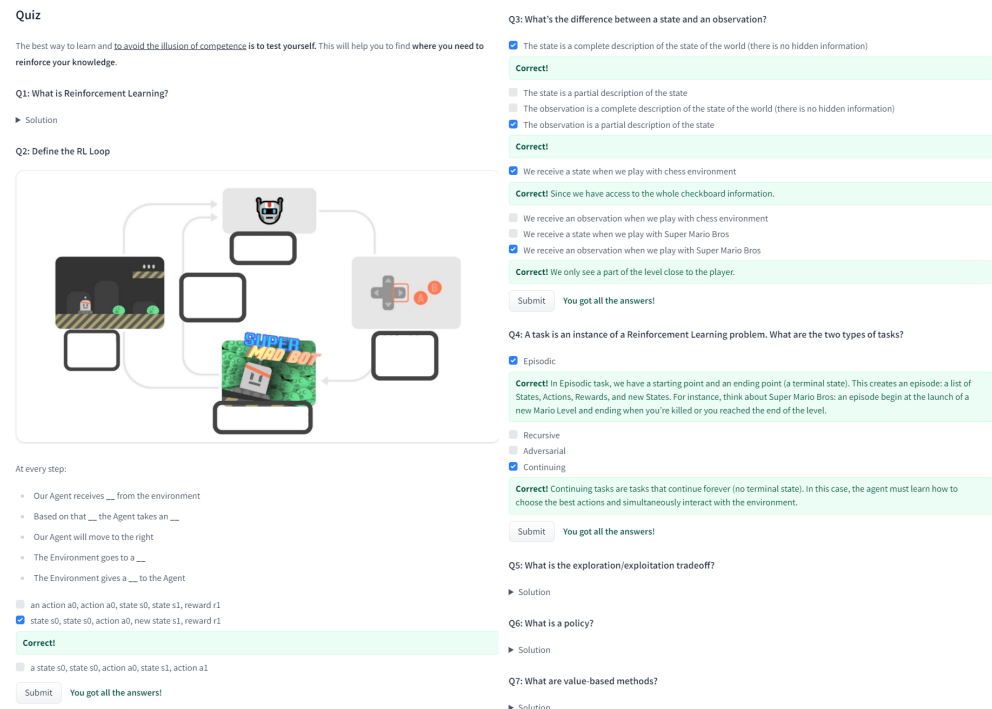
Solution

The finished units: 1, 2, 3, 4, 6, 8's codes could be check in 'code/Huggingface_Deep_RL'.

After running the implemented codes, several models from different units were pushed to the Hugging-face Hub, and the models are as follows:



The quizzes of each units are also finished, which could be seen as follows:



Mid-way Quiz

The best way to learn and to avoid the illusion of competence is to test yourself. This will help you to find where you need to reinforce your knowledge.

Q1: What are the two main approaches to find optimal policy?

☒ Policy-based methods

Correct! With Policy-Based methods, we train the policy directly to learn which action to take given a state.

☐ Random-based methods

☒ Value-based methods

Correct! With value-based methods, we train a value function to learn which state is more valuable and use this value function to take the action that leads to it.

☐ Evolution-strategies methods

You got all the answers!

Q2: What is the Bellman Equation?

► Solution

Q3: Define each part of the Bellman Equation

The Bellman Equation

$$V_{\pi}(s) = \mathbf{E}_{\pi}[R_{t+1} + \gamma * V_{\pi}(S_{t+1}) | S_t = s]$$

► Solution

Q4: What is the difference between Monte Carlo and Temporal Difference learning methods?

☒ With Monte Carlo methods, we update the value function from a complete episode

Correct!

☐ With Monte Carlo methods, we update the value function from a step

☐ With TD learning methods, we update the value function from a complete episode

☒ With TD learning methods, we update the value function from a step

Correct!

You got all the answers!

Q5: Define each part of Temporal Difference learning formula

TD Learning Approach:

Temporal Difference Learning: learning at each time step.

$$V(S_t) \leftarrow V(S_t) + \alpha[R_{t+1} + \gamma V(S_{t+1}) - V(S_t)]$$

► Solution

Q6: Define each part of Monte Carlo learning formula

Monte Carlo Approach:

Monte Carlo: waits until the end of the episode, then calculates Gt (return) and uses it as a target for its value or policy.

$$V(S_t) \leftarrow V(S_t) + \alpha[G_t - V(S_t)]$$

► Solution

Congrats on finishing this Quiz 🎉, if you missed some elements, take time to read again the previous sections to reinforce your knowledge.

Second Quiz

The best way to learn and to avoid the illusion of competence is to **test yourself**. This will help you to find **where you need to reinforce your knowledge**.

Q1: What is Q-Learning?

- ☒ The algorithm we use to train our Q-function

Correct!

- ☐ A value function
☐ An algorithm that determines the value of being at a particular state and taking a specific action at that state
☐ A table

Submit

You got all the answers!

Q2: What is a Q-table?

- ☐ An algorithm we use in Q-Learning
☒ Q-table is the internal memory of our agent

Correct!

- ☐ In Q-table each cell corresponds a state value

Submit

You got all the answers!

Q3: Why if we have an optimal Q-function Q^* we have an optimal policy?

► Solution

Q4: Can you explain what is Epsilon-Greedy Strategy?

► Solution

Q5: How do we update the Q value of a state, action pair?

$$\underline{Q(S_t, A_t)} \leftarrow \underline{Q(S_t, A_t)} + \alpha [\underline{R_{t+1}} + \gamma \underline{\max_a Q(S_{t+1}, a)} - \underline{Q(S_t, A_t)}]$$

► Solution

Q6: What's the difference between on-policy and off-policy

► Solution

Quiz

The best way to learn and to avoid the illusion of competence is to test yourself. This will help you to find where you need to reinforce your knowledge.

Q1: We mentioned Q Learning is a tabular method. What are tabular methods?

► Solution

Q2: Why can't we use a classical Q-Learning to solve an Atari Game?

- ☐ Atari environments are too fast for Q-Learning
- ☒ Atari environments have a big observation space. So creating an updating the Q-Table would not be efficient

Correct!

Submit

You got all the answers!

Q3: Why do we stack four frames together when we use frames as input in Deep Q-Learning?

► Solution

Q4: What are the two phases of Deep Q-Learning?

- ☒ Sampling

Correct! We perform actions and store the observed experiences tuples in a replay memory.

- ☐ Shuffling
- ☐ Reranking
- ☒ Training

Correct! We select the small batch of tuple randomly and learn from it using a gradient descent update step.

Submit

You got all the answers!

Q5: Why do we create a replay memory in Deep Q-Learning?

► Solution

Q6: How do we use Double Deep Q-Learning?

► Solution

Congrats on finishing this Quiz 🎉, if you missed some elements, take time to read again the chapter to reinforce (🧐) your

Quiz

The best way to learn and to avoid the illusion of competence is to test yourself. This will help you to find where you need to reinforce your knowledge.

Q1: What are the advantages of policy-gradient over value-based methods? (Check all that apply)

☒ Policy-gradient methods can learn a stochastic policy

Correct!

☒ Policy-gradient methods are more effective in high-dimensional action spaces and continuous actions spaces

Correct!

☐ Policy-gradient converges most of the time on a global maximum.

Submit

You got all the answers!

Q2: What is the Policy Gradient Theorem?

► Solution

Q3: What's the difference between policy-based methods and policy-gradient methods? (Check all that apply)

☐ Policy-based methods are a subset of policy-gradient methods.

☒ Policy-gradient methods are a subset of policy-based methods.

Correct!

☒ In Policy-based methods, we can optimize the parameter θ **indirectly** by maximizing the local approximation of the objective function with techniques like hill climbing, simulated annealing, or evolution strategies.

Correct!

☒ In Policy-gradient methods, we optimize the parameter θ **directly** by performing the gradient ascent on the performance of the objective function.

Correct!

Submit

You got all the answers!

Q4: Why do we use gradient ascent instead of gradient descent to optimize $J(\theta)$?

☐ We want to minimize $J(\theta)$ and gradient ascent gives us the direction of the steepest increase of $J(\theta)$

☒ We want to maximize $J(\theta)$ and gradient ascent gives us the direction of the steepest increase of $J(\theta)$

Correct!

Submit

You got all the answers!

Congrats on finishing this Quiz 🎉, if you missed some elements, take time to read the chapter again to reinforce 😊 your knowledge.

Quiz

The best way to learn and to avoid the illusion of competence is to test yourself. This will help you to find where you need to reinforce your knowledge.

Q1: Which of the following interpretations of bias-variance tradeoff is the most accurate in the field of Reinforcement Learning?

- ☐ The bias-variance tradeoff reflects how my model is able to generalize the knowledge to previously tagged data we give to the model during training time.
- ☒ The bias-variance tradeoff reflects how well the reinforcement signal reflects the true reward the agent should get from the environment

Correct!

Submit

You got all the answers!

Q2: Which of the following statements are true, when talking about models with bias and/or variance in RL?

- ☒ An unbiased reward signal returns rewards similar to the real / expected ones from the environment

Correct!

- ☐ A biased reward signal returns rewards similar to the real / expected ones from the environment
- ☒ A reward signal with high variance has much noise in it and gets affected by, for example, stochastic (non constant) elements in the environment

Correct!

- ☐ A reward signal with low variance has much noise in it and gets affected by, for example, stochastic (non constant) elements in the environment

Submit

You got all the answers!

Q3: Which of the following statements are true about Monte Carlo method?

- ☒ It's a sampling mechanism, which means we don't analyze all the possible states, but a sample of those

Correct!

- ☐ It's very resistant to stochasticity (random elements in the trajectory)
- ☒ To reduce the impact of stochastic elements in Monte Carlo, we take 'n' strategies and average them, reducing their individual impact

Correct!

Submit

You got all the answers!

Q4: How would you describe, with your own words, the Actor-Critic Method (A2C)?

► Solution

Q5: Which of the following statements are true about the Actor-Critic Method?

- ☐ The Critic does not learn any function during the training process
- ☒ The Actor learns a policy function, while the Critic learns a value function

Correct!

- ☒ It adds resistance to stochasticity and reduces high variance

Correct!

Submit

You got all the answers!